

HW7: Convolutional Neural Networks


Due Mar 30, 2021 by 2:30pm **Points** 100 **Submitting** a file upload
Available Mar 23, 2021 at 2:30pm - Mar 30, 2021 at 2:30pm 7 days

This assignment was locked Mar 30, 2021 at 2:30pm.

Assignment Goals

- Compare creating a CNN to a basic neural network as in HW6
- Design your own deep network for scene recognition

Summary

Your implementation in this assignment might take one or two hours to run. **We highly recommend to start working on this assignment early!** In this homework, we will explore building deep neural networks, including Convolutional Neural Networks (CNNs), using PyTorch. Helper code is provided in this assignment and can be downloaded from this zip file [code_hw7-2.zip](https://canvas.wisc.edu/courses/230450/files/19086157/download?download_frd=1)  [\(https://canvas.wisc.edu/courses/230450/files/19086157/download?download_frd=1\)](https://canvas.wisc.edu/courses/230450/files/19086157/download?download_frd=1) .

In this HW, you are still expected to use the Pytorch backend and virtual environment for programming where you can find the tutorials in [HW6](#).

Design a CNN model for MiniPlaces Dataset

In this part, you will design a simple CNN model along with your own convolutional network for a more realistic dataset -- [MiniPlaces](https://github.com/CSAILVision/miniplaces) [_ \(https://github.com/CSAILVision/miniplaces\)](https://github.com/CSAILVision/miniplaces) , again using PyTorch.

Dataset

MiniPlaces is a scene recognition dataset developed by MIT. This dataset has 120K images from 100 scene categories. The categories are mutually exclusive. The dataset is split into 100K images for training, 10K images for validation, and 10K for testing.

Helper Code

We provide helper functions in **train_miniplaces.py** and **dataloader.py**, and skeleton code in **student_code.py**. See the comments in these files for more implementation details.

The original image resolution for images in MiniPlaces is 128x128. To make the training feasible, our data loader reduces the image resolution to 32x32. You can always assume this input resolution. Our

data loader will also download the full dataset the first time you run `train_miniplaces.py`

Before the training procedure, we define the dataloader, model, optimizer, image transform and criterion. We execute the training and testing in function **train_model** and **test_model**, which is similar to what we have for HW6.

Part I Creating LeNet-5

In this part, you have to implement a classic CNN model, called LeNet-5 in Pytorch for the MiniPlaces dataset. We use the following layers in this order:

1. One convolutional layer with the number of output channels to be 6, kernel size to be 5, stride to be 1, padding to be 0 followed by a relu activation layer and then a 2D max pooling layer (kernel size to be 2, padding to be 0 and stride to be 2).
2. One convolutional layer with the number of output channels to be 16, kernel size to be 5, stride to be 1, padding to be 0 followed by a relu activation layer and then a 2D max pooling layer (kernel size to be 2, padding to be 0 and stride to be 2).
3. A Flatten layer to convert the 3D tensor to a 1D tensor.
4. A Linear layer with output dimension to be 256, followed by a relu activation function.
5. A Linear layer with output dimension to be 128, followed by a relu activation function.
6. A Linear layer with output dimension to be the number of classes (in our case, 100).

You have to fill in the **LeNet** class in **student_code.py**. You are expected to create the model following [this tutorial](https://pytorch.org/tutorials/beginner/examples_nn/two_layer_net_module.html) (https://pytorch.org/tutorials/beginner/examples_nn/two_layer_net_module.html), which is different from using `nn.Sequential()` in the last HW.

In addition, given a batch of inputs with shape `[N, C, W, H]`, where `N` is the batch size, `C` is the input channel and `W, H` are the width and height of the image (both 32 in our case), you are expected to return both the output of the model along with the shape of the intermediate outputs for the above 6 stages. The shape should be a dictionary with the keys to be 1,2,3,4,5,6 (integers) denoting each stage, and the corresponding value to be a list that denotes the shape of the intermediate outputs.

Hint: The expected model has the following form:

```
class LeNet(nn.Module):  
  
    def __init__(self, input_shape=(32, 32), num_classes=100):  
  
        super(LeNet, self).__init__()  
  
        # certain definitions  
  
    def forward(self, x):
```

```
shape_dict = {}
```

```
# certain operations
```

```
return out, shape_dict
```

Shape_dict should have the following form: {1: [a,b,c,d], 2:[e,f,g,h], ..., 6: [x,y]}

The linear layer and the convolutional layer have bias terms.

Part II Count the number of trainable parameters of LeNet-5

In this part, you are expected to return the number of trainable parameters of your created LeNet model in the previous part. You have to fill in the function **count_model_params** in **student_code.py**.

The function output is in the unit of Million(1e6). **Please do not use any external library which directly calculates the number of parameters (other libraries, such as NumPy can be used)!**

Hint: You can use the `model.named_parameters()` to gain the name and the corresponding parameters of a model. Please do not do any rounding to the result.

Part III Training LeNet-5 under different configurations

Based on the LeNet-5 model created in the previous parts, in this section, you are expected to train the LeNet-5 model under different configurations.

You will use similar implementations of **train_model** and **test_model** as you did for [HW6](#) (which we provide in `student_code.py`). When you run `train_miniplaces.py`, the python script will save two files in the "outputs" folder.

- **checkpoint.pth.tar** is the model checkpoint at the latest epoch.
- **model_best.pth.tar** is the model weights that has highest accuracy on the validation set.

Our code supports resuming from a previous checkpoint, such that you can pause the training and resume later. This can be achieved by running

```
python train_miniplaces.py --resume ./outputs/checkpoint.pth.tar
```

After training, we provide **eval_miniplaces.py** to help you evaluate the model on the validation set and also help you in timing your model. This script will grab a pre-trained model and evaluate it on the validation set of 10K images. For example, you can run

```
python eval_miniplaces.py --load ./outputs/model_best.pth.tar
```

The output shows the validation accuracy and also the model evaluation time in seconds (see an example below).

```
=> Loading from cached file ../data/miniplaces/cached_val.pkl ...
```

```
=> loading checkpoint '../outputs/model_best.pth.tar'
```

```
=> loaded checkpoint '../outputs/model_best.pth.tar' (epoch x)
```

```
Evaluating the model ...
```

```
[Test set] Epoch: xx, Accuracy: xx.xx%
```

```
Evaluation took 2.26 sec
```

You can run this script a few times to see the average runtime of your model.

Please train the model under the following different configurations:

1. The default configuration provided in the code, which means you do not have to make modifications.
2. Set the batch size to 8, the remaining configurations are kept the same.
3. Set the batch size to 16, the remaining configurations are kept the same.
4. Set the learning rate to 0.05, the remaining configurations are kept the same.
5. Set the learning rate to 0.01, the remaining configurations are kept the same.
6. Set the epochs to 20, the remaining configurations are kept the same.
7. Set the epochs to 5, the remaining configurations are kept the same.

After training, you are expected to get the validation accuracy for each configuration using the best model (**model_best.pth.tar**), then save these accuracies into a **results.txt** file, where the accuracy of each configuration is placed in one line in order. Your .txt file will end up looking like this:

```
18.00
```

```
17.39
```

```
12.31
```

```
...
```

The exact accuracy may not align well with your results. They are just for illustration purposes.

You have to submit the **results.txt** file together with your **student_code.py**.

Optional: Profiling Your Model

You might find that the training or evaluation of your model is a bit slower than expected. Fortunately,

PyTorch has its own profiling tool. Here is a quick [tutorial](https://pytorch.org/tutorials/recipes/recipes/profiler.html) [_ \(https://pytorch.org/tutorials/recipes/recipes/profiler.html\)](https://pytorch.org/tutorials/recipes/recipes/profiler.html) of using PyTorch profiler. You can easily inject the profiler into `train_miniplaces.py` to inspect the runtime and memory consumption of different parts of your model. A general principle is that a deep (many layers) and wide (many feature channels) network will train much slower. It is your design choice to balance between the efficiency and the accuracy.

Optional: Training on CSL

We recommend training the model on your local machine or CSL machine and time your model (using `eval_miniplaces.py`) on CSL machines. If you decide to train the model on a CSL machine, you will need to find a way to allow your remote session to remain active when you are disconnected from CSL. In this case, we recommend using `tmux`, a terminal multiplexer for Unix-like systems. `tmux` is already installed on CSL. To use `tmux`, simply type **tmux** in the terminal. Now you can run your code in a `tmux` session. And the session will remain active even if you are disconnected.

- If you want to detach a `tmux` session without closing it, press "ctrl + b" then "d" (detach) within a `tmux` session. This will exist to the terminal while keeping the session active. Later you can re-attach the session.
- If you want to enter an active `tmux` session, type "tmux a" to attach to the last session in the terminal (outside of `tmux`).
- If you want to close a `tmux` session, press "ctrl + b" then "x" (exit) within a `tmux` session. You won't be able to enter this session again. Please make sure that you close your `tmux` sessions after this assignment.
- See here for a [brief tutorial](https://linuxize.com/post/getting-started-with-tmux/) [_ \(https://linuxize.com/post/getting-started-with-tmux/\)](https://linuxize.com/post/getting-started-with-tmux/) on the powerful tool of `tmux`.

Deliverable

You will need to submit **student_code.py** together with your **results.txt** for Part III. The validation accuracy for different configurations must be in **.txt** format and named as **results.txt**.

Submission Notes

You will need to submit a single zip file named **hw7_<netid>.zip**, where you replace `<netid>` with your netID (your wisc.edu login). This file will include `student_code.py` and the validation accuracy file with the proper name as mentioned above. Do not submit a Jupyter notebook `.ipynb` file.

Be sure to **remove all debugging output** before submission. Failure to remove debugging output will be **penalized (10pts)**.

If a regrading request isn't justifiable (the initial grade is correct and clear, subject to the

instructors' judgment), the request for regrading will be penalized (10 pts).

This assignment is due on March 30, 2:30 PM. We highly recommend to start early. It is preferable to first submit a version well before the deadline (at least one hour before) and check the content/format of the submission to make sure it's the right version. Then, later update the submission until the deadline if needed.

Scoring Rubric			
Criteria	Ratings		Pts
Correctly create the LeNet model with the right output shape.	30 pts Full Marks	0 pts No Marks	30 pts
Correctly calculate the number of parameters of LeNet.	20 pts Full Marks	0 pts No Marks	20 pts
Correctly train the model under different configurations and report the right accuracy.	50 pts Full Marks	0 pts No Marks	50 pts
			Total Points: 100